

# *SQL*

“Structured Query Language”

# *SQL*

The swiss army knife for data



guest talk by DC Posch

<https://dcpos.ch>

# *SQL Introduction*

Pros:

**Universal. Declarative. Timeless,  
never goes out of fashion.**

# *SQL Introduction*

Pros:

Universal. Declarative. Timeless,  
never goes out of fashion.

**A powerful tool both for production  
use and for data analysis.**

# *SQL Introduction*

## Plan

- SQL Crash Course + challenge
- Core Concepts + challenge
- Advanced SQL
- Conclusion

# *SQL Crash Course*

# *SQL Crash Course*

**SELECT** to read (and transform, and summarize, and analyze) data.

**INSERT, UPDATE, DELETE** to write data.

*Reading data*



# *SQL Crash Course*

```
SELECT *
```

```
FROM plays;
```

## *SQL Crash Course*

```
SELECT play_id, type, player,  
        assist, points, description  
FROM plays;
```

# *SQL Crash Course*

**Try it! Please fork it:**

<https://www.db-fiddle.com/f/tzUp9abziH3mXxvhWoeKMm/0>

*Filtering*

# *SQL Crash Course*

```
SELECT *
```

```
FROM plays
```

```
WHERE player='Lebron James'
```

# *SQL Crash Course*

```
SELECT *
```

```
FROM plays
```

```
WHERE player='Lebron James'
```

```
AND event_type in ('shot', 'miss')
```

*Grouping*

## *SQL Crash Course*

```
SELECT team, player, points  
FROM plays;
```



## *SQL Crash Course*

```
SELECT team, player, sum(points)
FROM plays
GROUP BY team, player
ORDER BY sum(points) desc;
```

## *SQL Crash Course*

```
SELECT player, event_type, count(*)  
FROM plays  
WHERE event_type in ('shot', 'miss')  
GROUP BY player, event_type;
```

# *ChalLenge*

player	shotsTaken	shotsMade
Anthony Davis	24	10
	...	

*Questions so far?*

# *SQL Concepts*

# *SQL Concepts*

Relation (n): fancy word for row

Foreign keys and constraints

Join: combining data from two tables

Transaction: updating data atomically

ACID: atomicity, consistency,  
isolation, and durability

*Join*

## *SQL Concepts*

```
SELECT *
```

```
FROM order_cart_items ci
```

```
INNER JOIN items i ON i.id=ci.item_i;
```



## *SQL Concepts*

```
SELECT brand, sum(quantity)  
FROM order_cart_items ci  
INNER JOIN items i ON i.id=ci.item_id  
GROUP BY i.brand;
```

## *SQL Concepts*

```
SELECT brand, sum(quantity)
FROM order_cart_items ci
JOIN items i ON i.id=ci.item_id
GROUP BY i.brand;
```

# *SQL Concepts*

**Quick about about minimizing roundtrips**

## *SQL Concepts*

```
insert into order_cart_items  
(order_id,item_id,quantity)  
  values (1, 1, 3),(1, 3, 2),(1, 4, 1),  
  (2, 1, 1),(2, 4, 1);
```

# *SQL Concepts*

Try it, this time with joins!

Please fork:

<https://www.db-fiddle.com/f/pTChybKEAvb6krkJdo3sb3/3>

# *Left join*

*For when you want all the rows in A,  
even when there's no match over in B*

## *SQL Concepts*

```
SELECT name, sum(quantity)
FROM items i
JOIN order_cart_items ci
ON i.id=ci.item_id
GROUP BY name;
```

# *ChalLenge*

*Fix that query to show the missing item.*



# *Advanced SQL*

*Transactions*

## *SQL Advanced: Transactions*

```
BEGIN TRANSACTION
```

```
UPDATE accounts SET bal=bal-amount  
WHERE id=100;
```

```
UPDATE accounts SET bal=bal+amount  
WHERE id=200;
```

```
COMMIT
```

## *SQL Advanced: Transactions*

**The core intuition for transactions: a transaction either happens all the way, or not at all. "Atomicity"**

# *SQL Advanced: Transactions*

**ACID** = Atomicity +

**Isolation** = you'll never see a "half-way" finished transaction.

**Consistency** = constraints still hold

**Durability** = what's done is done

# *Normalization*

## *SQL Advanced: Normalization*

**Every bit of information should have a single source of truth. Don't paste the same stuff into two different tables.**

*Window functions*



## *SQL Advanced: Window functions*

Window functions let you do rankings, running totals, and similar. They allow a result row from a query to depend on the contents of *\*other\** result rows.

## *SQL Advanced: Window Functions*

```
SELECT ...,
```

```
    RANK(points) OVER(PARTITION BY team)
```

```
FROM plays;
```

*JSON*

## *SQL Advanced: JSON*

- Lets you mix document-style with relational DB.
- Saves you a ton of tables.
- Makes it easier to atomically update a whole complex object.
- Rule of thumb: no ID references inside of JSON. Just plain data.

## *SQL Advanced: JSON*

```
SELECT
```

```
    id, docId
```

```
    slideJson
```

```
FROM slides;
```

## *SQL Advanced: JSON*

**SELECT**

**slideJson->'\$.bgColor' as col,**

**count(\*)**

**FROM slides**

**GROUP BY bgColor;**

*Quick demo*

# *Questions*

dcposch@dcpos.ch